



Tensors

```
import torch
import numpy as np
```

```
data = [[1, 2],[3, 4]]
x_data = torch.tensor(data)
```

```
np_array = np.array(data)
x_np = torch.from_numpy(np_array)
```

```
shape = (2,3,)
rand_tensor = torch.rand(shape)
ones_tensor = torch.ones(shape)
zeros_tensor = torch.zeros(shape)
```

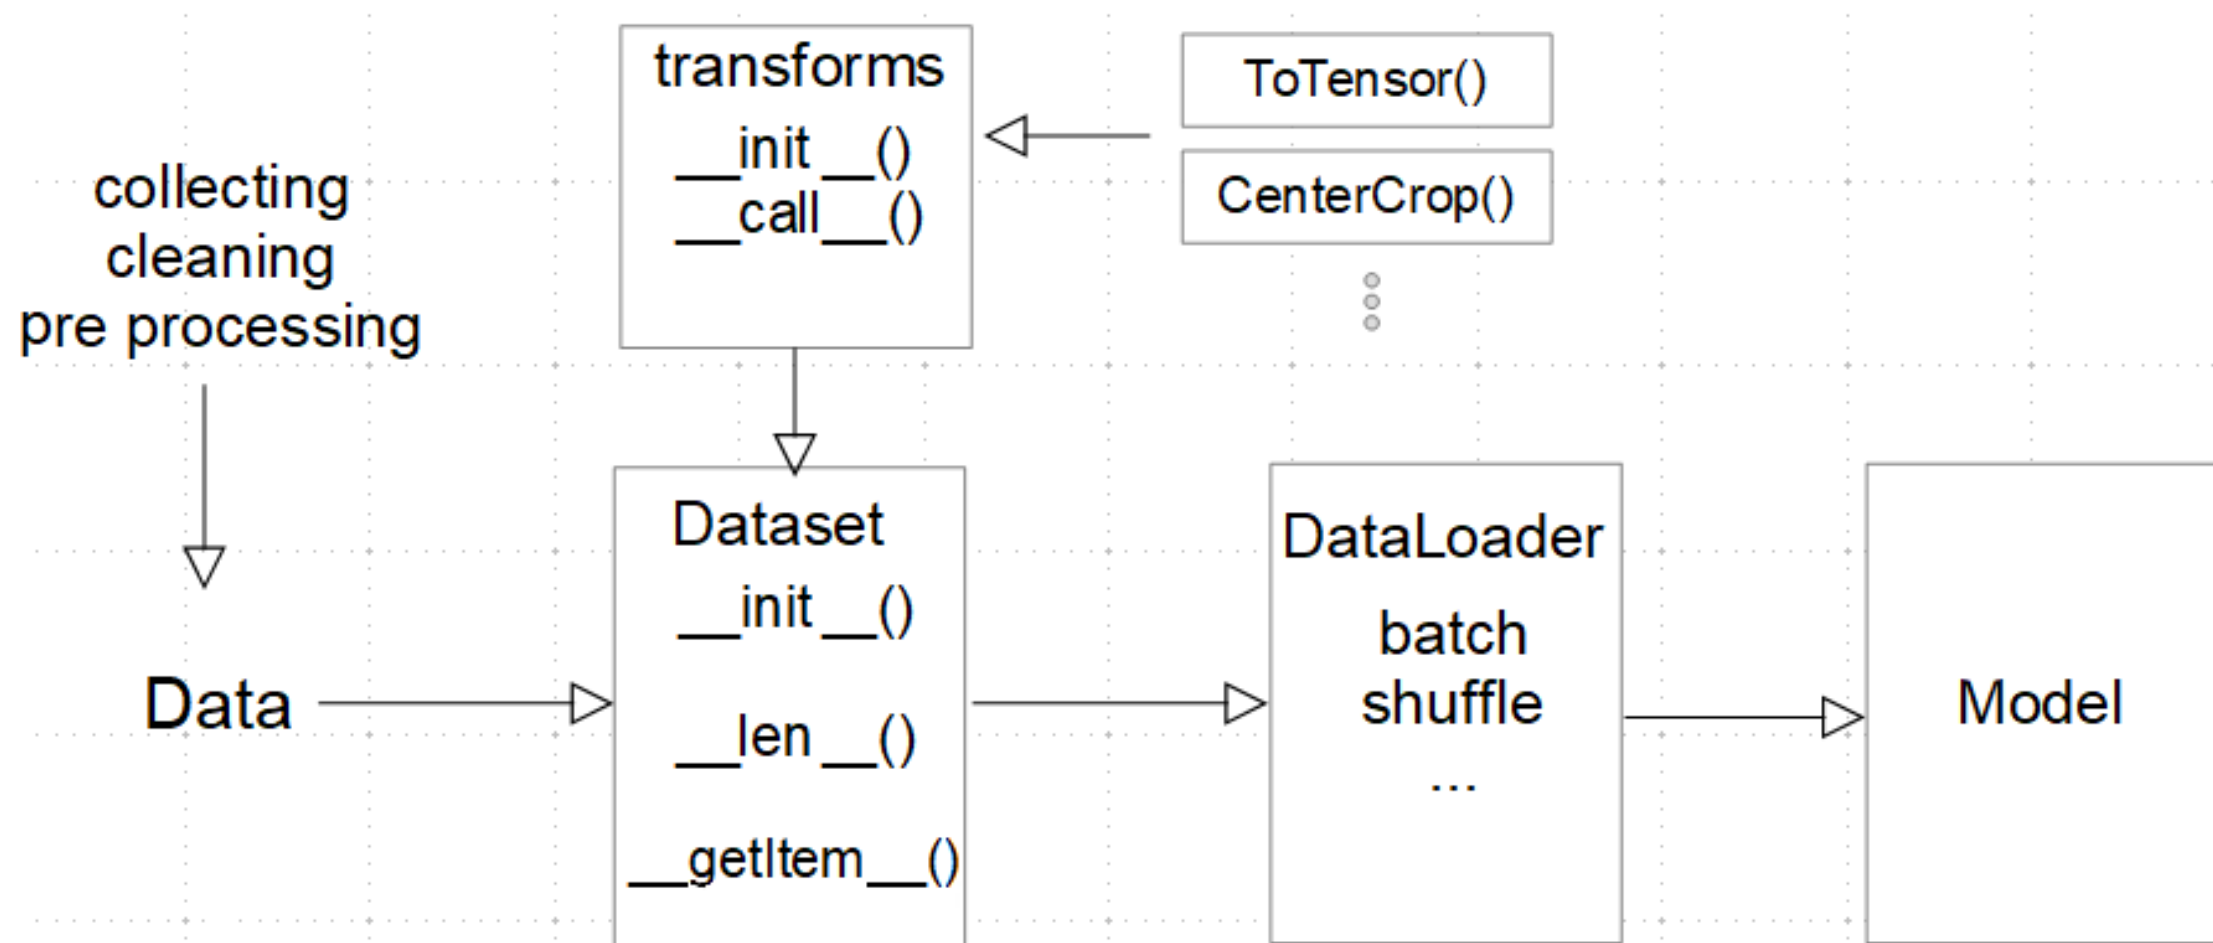
```
y1 = tensor @ tensor.T
```

```
z1 = tensor * tensor
```

Tensors

```
# We move our tensor to the GPU if available  
if torch.cuda.is_available():  
    tensor = tensor.to("cuda")
```

Datasets & Data loaders



Test dataset

```
import torch
from torch.utils.data import Dataset
from torchvision import datasets
from torchvision.transforms import ToTensor
import matplotlib.pyplot as plt
```

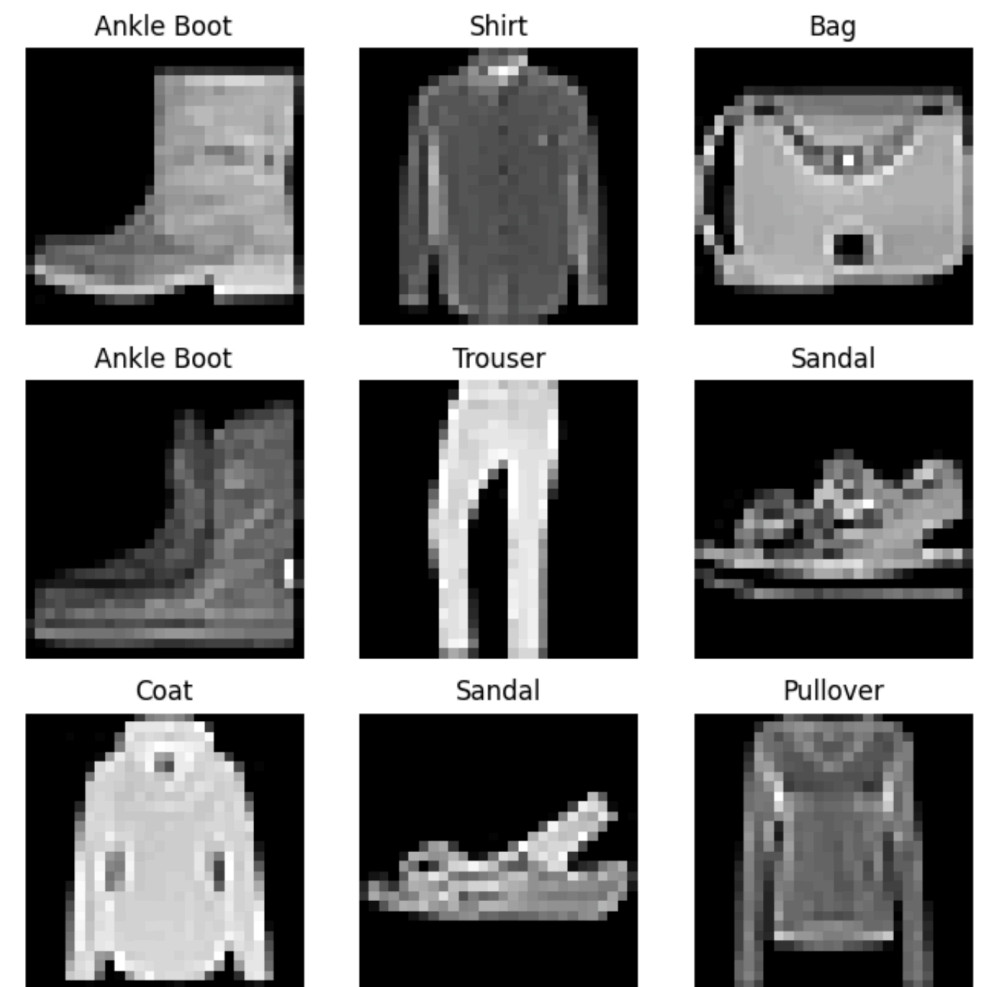
```
training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor()
)
```

```
test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor()
)
```

Fashion MNIST

```
labels_map = {
    0: "T-Shirt",
    1: "Trouser",
    2: "Pullover",
    3: "Dress",
    4: "Coat",
    5: "Sandal",
    6: "Shirt",
    7: "Sneaker",
    8: "Bag",
    9: "Ankle Boot",
}

figure = plt.figure(figsize=(8, 8))
cols, rows = 3, 3
for i in range(1, cols * rows + 1):
    sample_idx = torch.randint(len(training_data), size=(1,)).item()
    img, label = training_data[sample_idx]
    figure.add_subplot(rows, cols, i)
    plt.title(labels_map[label])
    plt.axis("off")
    plt.imshow(img.squeeze(), cmap="gray")
plt.show()
```



Custom dataset

```
import os
import pandas as pd
from torchvision.io import read_image

class CustomImageDataset(Dataset):
    def __init__(self, annotations_file, img_dir, transform=None, target_transform=None):
        self.img_labels = pd.read_csv(annotations_file)
        self.img_dir = img_dir
        self.transform = transform
        self.target_transform = target_transform

    def __len__(self):
        return len(self.img_labels)

    def __getitem__(self, idx):
        img_path = os.path.join(self.img_dir, self.img_labels.iloc[idx, 0])
        image = read_image(img_path)
        label = self.img_labels.iloc[idx, 1]
        if self.transform:
            image = self.transform(image)
        if self.target_transform:
            label = self.target_transform(label)
        return image, label
```

Data loaders

```
from torch.utils.data import DataLoader
```

```
train_dataloader = DataLoader(training_data, batch_size=64, shuffle=True)
```

```
test_dataloader = DataLoader(test_data, batch_size=64, shuffle=True)
```

```
# Display image and label.
```

```
train_features, train_labels = next(iter(train_dataloader))
```

```
print(f"Feature batch shape: {train_features.size()}")
```

```
print(f"Labels batch shape: {train_labels.size()}")
```

```
img = train_features[0].squeeze()
```

```
label = train_labels[0]
```

```
plt.imshow(img, cmap="gray")
```

```
plt.show()
```

```
print(f"Label: {label}")
```


Building neural network

```
import os
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
```

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```

Building neural network

```
model = NeuralNetwork().to(device)
print(model)
```

```
NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=784, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=512, bias=True)
    (3): ReLU()
    (4): Linear(in_features=512, out_features=10, bias=True)
  )
)
```

```
device = (
  "cuda"
  if torch.cuda.is_available()
  else "mps"
  if torch.backends.mps.is_available()
  else "cpu"
)
```

Automatic Differentiation

```
import torch
```

```
x = torch.ones(5)  # input tensor
```


```
y = torch.zeros(3)  # expected output
```

```
w = torch.randn(5, 3, requires_grad=True)
```

```
b = torch.randn(3, requires_grad=True)
```

```
z = torch.matmul(x, w)+b
```

```
loss = torch.nn.functional.binary_cross_entropy_with_logits(z, y)
```



```
print(f"Gradient function for z = {z.grad_fn}")
```

```
print(f"Gradient function for loss = {loss.grad_fn}")
```


```
Gradient function for z = <AddBackward0 object at 0x7ff0a3df3010>
```

```
Gradient function for loss = <BinaryCrossEntropyWithLogitsBackward0 object at  
0x7ff0a2e813f0>
```

Automatic Differentiation

```
loss.backward()  
print(w.grad)  
print(b.grad)
```

```
tensor([[0.3313, 0.0626, 0.2530],  
        [0.3313, 0.0626, 0.2530],  
        [0.3313, 0.0626, 0.2530],  
        [0.3313, 0.0626, 0.2530],  
        [0.3313, 0.0626, 0.2530]])  
tensor([0.3313, 0.0626, 0.2530])
```



Training the neural network

```
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor

training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor()
)

test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor()
)

train_dataloader = DataLoader(training_data, batch_size=64)
test_dataloader = DataLoader(test_data, batch_size=64)
```

Training the neural network

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

model = NeuralNetwork()
```

Training the neural network

```
learning_rate = 1e-3
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

epochs = 10
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_dataloader, model, loss_fn, optimizer)
    test_loop(test_dataloader, model, loss_fn)
print("Done!")
```

Training the neural network

```
def train_loop(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    # Set the model to training mode - important for batch normalization and dropout layers
    # Unnecessary in this situation but added for best practices
    model.train()
    for batch, (X, y) in enumerate(dataloader):
        # Compute prediction and loss
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

    if batch % 100 == 0:
        loss, current = loss.item(), batch * batch_size + len(X)
        print(f"loss: {loss:>7f}  [{current:>5d}/{size:>5d}"])
```


Training the neural network

```
def test_loop(dataloader, model, loss_fn):  
    # Set the model to evaluation mode - important for batch normalization and dropout layers  
    # Unnecessary in this situation but added for best practices  
    model.eval()  
    size = len(dataloader.dataset)  
    num_batches = len(dataloader)  
    test_loss, correct = 0, 0  
  
    # Evaluating the model with torch.no_grad() ensures that no gradients are computed during  
    test mode  
    # also serves to reduce unnecessary gradient computations and memory usage for tensors with  
    requires_grad=True  
    with torch.no_grad():  
        for X, y in dataloader:  
            pred = model(X)  
            test_loss += loss_fn(pred, y).item()  
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()  
  
    test_loss /= num_batches  
    correct /= size  
    print(f"Test Error: \n Accuracy: {(100*correct)/size}>0.1f}%, Avg loss: {test_loss:>8f} \n")
```



PyTorch